

VIET NAM NATIONAL UNIVERSITY HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



SEMESTER 241 - CO3071

LAB 5 REPORT

DISTRIBUTED SYSTEMS

Instructor: PROF. Nguyễn Mạnh Thìn

Tạ Gia Khang	- 2152642
Nguyễn Quang Thiện	- 2152994

HO CHI MINH CITY, DECEMBER 2024



Contents

1	Homework	3
1.1	Configuration	3
1.2	Previous LAB's Output	8
1.3	Loading Data	9
1.4	Bring data to compute (Centralized Computing)	11
1.5	Bring compute to data (Distributed Computing)	13
1.6	Store Output	15
1.7	Demonstration	16
1.7.1	Centralized Computing Method	16
1.7.2	Distributed Computing Method	17

1 Homework

1.1 Configuration

Using Docker and generating a docker-compose.yml file to create 3 brokers for **Kafka Fog**, **Spark Master**, **2 Spark Workers**, a **Hadoop Namenode**, and **3 Datanodes**.

```
1  version: '2'
2  services:
3    # kafka-fog
4    kafka1:
5      image: apache/kafka:3.7.0
6      hostname: kafka1
7      container_name: kafka1
8      ports:
9        - 29092:9092
10     environment:
11       KAFKA_NODE_ID: 1
12       KAFKA_PROCESS_ROLES: 'broker,controller'
13       KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
14         'CONTROLLER:PLAINTEXT,PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT'
15       KAFKA_CONTROLLER_QUORUM_VOTERS: '1@kafka1:9093,2@kafka2:9093,3@kafka3:9093'
16       KAFKA_LISTENERS:
17         'PLAINTEXT://:29092,CONTROLLER://:9093,PLAINTEXT_HOST://:9092'
18       KAFKA_INTER_BROKER_LISTENER_NAME: 'PLAINTEXT'
19       KAFKA_ADVERTISED_LISTENERS:
20         PLAINTEXT://kafka1:29092,PLAINTEXT_HOST://localhost:29092
21       KAFKA_CONTROLLER_LISTENER_NAMES: 'CONTROLLER'
22       CLUSTER_ID: '2UKPPoqNTPezeassrAogQw'
23       KAFKA_OFFSETS_TOPIC_NUM_PARTITIONS: 1
24       KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
25       KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1
26       KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1
27       KAFKA_LOG_DIRS: '/tmp/kraft-combined-logs'
28     networks:
29       - lab5_ds_kafka_network
30
31   kafka2:
32     image: apache/kafka:3.7.0
33     hostname: kafka2
34     container_name: kafka2
35     ports:
36       - 39092:9092
37     environment:
38       KAFKA_NODE_ID: 2
39       KAFKA_PROCESS_ROLES: 'broker,controller'
```

```
38 KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
39 'CONTROLLER:PLAINTEXT,PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT'
40 KAFKA_CONTROLLER_QUORUM_VOTERS: '1@kafka1:9093,2@kafka2:9093,3@kafka3:9093'
41 KAFKA_LISTENERS:
42 'PLAINTEXT://:39092,CONTROLLER://:9093,PLAINTEXT_HOST://:9092'
43 KAFKA_INTER_BROKER_LISTENER_NAME: 'PLAINTEXT'
44 KAFKA_ADVERTISED_LISTENERS:
45 PLAINTEXT://kafka2:39092,PLAINTEXT_HOST://localhost:39092
46 KAFKA_CONTROLLER_LISTENER_NAMES: 'CONTROLLER'
47 CLUSTER_ID: '2UKPPoqNTPezeassrAogQw'
48 KAFKA_OFFSETS_TOPIC_NUM_PARTITIONS: 1
49 KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
50 KAFKA_GROUP_INITIAL_REBALANCE_DELAY_MS: 0
51 KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1
52 KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1
53 KAFKA_LOG_DIRS: '/tmp/kraft-combined-logs'
54 networks:
55 - lab5_ds_kafka_network
56
57 kafka3:
58 image: apache/kafka:3.7.0
59 hostname: kafka3
60 container_name: kafka3
61 ports:
62 - 49092:9092
63 environment:
64 KAFKA_NODE_ID: 3
65 KAFKA_PROCESS_ROLES: 'broker,controller'
66 KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
67 'CONTROLLER:PLAINTEXT,PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT'
68 KAFKA_CONTROLLER_QUORUM_VOTERS: '1@kafka1:9093,2@kafka2:9093,3@kafka3:9093'
69 KAFKA_LISTENERS:
70 'PLAINTEXT://:49092,CONTROLLER://:9093,PLAINTEXT_HOST://:9092'
71 KAFKA_INTER_BROKER_LISTENER_NAME: 'PLAINTEXT'
72 KAFKA_ADVERTISED_LISTENERS:
73 PLAINTEXT://kafka3:49092,PLAINTEXT_HOST://localhost:49092
74 KAFKA_CONTROLLER_LISTENER_NAMES: 'CONTROLLER'
75 CLUSTER_ID: '2UKPPoqNTPezeassrAogQw'
76 KAFKA_OFFSETS_TOPIC_NUM_PARTITIONS: 1
77 KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
78 KAFKA_GROUP_INITIAL_REBALANCE_DELAY_MS: 0
79 KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1
80 KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1
81 KAFKA_LOG_DIRS: '/tmp/kraft-combined-logs'
82 networks:
83 - lab5_ds_kafka_network
```

```
79 kafka-ui:
80     container_name: kafka-ui
81     image: provectuslabs/kafka-ui:latest
82     ports:
83         - "8080:8080"
84     environment:
85         DYNAMIC_CONFIG_ENABLED: 'true'
86     depends_on:
87         - kafka1
88         - kafka2
89         - kafka3
90     networks:
91         - lab5_ds_kafka_network
92
93 spark-master:
94     image: bitnami/spark:latest
95     container_name: spark-master
96     environment:
97         - SPARK_MODE=master
98         - SPARK_RPC_AUTHENTICATION_ENABLED=no
99         - SPARK_RPC_ENCRYPTION_ENABLED=no
100        - SPARK_LOCAL_STORAGE_ENCRYPTION_ENABLED=no
101        - SPARK_SSL_ENABLED=no
102     ports:
103         - "7077:7077" # Spark Master port
104     networks:
105         - lab5_ds_kafka_network
106     volumes:
107         - ./pyspark:/opt/spark/python
108
109 spark-worker1:
110     image: bitnami/spark:latest
111     container_name: spark-worker1
112     environment:
113         - SPARK_MODE=worker
114         - SPARK_MASTER_URL=spark://spark-master:7077
115     depends_on:
116         - spark-master
117     networks:
118         - lab5_ds_kafka_network
119
120 spark-worker2:
121     image: bitnami/spark:latest
122     container_name: spark-worker2
123     environment:
124         - SPARK_MODE=worker
125         - SPARK_MASTER_URL=spark://spark-master:7077
```

```
126     depends_on:
127         - spark-master
128     networks:
129         - lab5_ds_kafka_network
130
131     namenode:
132         image: bde2020/hadoop-namenode:latest
133         container_name: namenode
134         environment:
135             - CLUSTER_NAME=hadoop_cluster
136             - CORE_CONF_fs_defaultFS=hdfs://namenode:9000
137         ports:
138             - "9000:9000" # HDFS namenode RPC port
139             - "9870:9870" # HDFS namenode Web UI
140         networks:
141             - lab5_ds_kafka_network
142         volumes:
143             - ./namenode:/hadoop/dfs/name
144
145     datanode:
146         image: bde2020/hadoop-datanode:latest
147         container_name: datanode
148         environment:
149             - CORE_CONF_fs_defaultFS=hdfs://namenode:9000
150         depends_on:
151             - namenode
152         networks:
153             - lab5_ds_kafka_network
154         volumes:
155             - ./datanode:/hadoop/dfs/data
156
157     datanode2:
158         image: bde2020/hadoop-datanode:latest
159         container_name: datanode2
160         environment:
161             - CORE_CONF_fs_defaultFS=hdfs://namenode:9000
162         depends_on:
163             - namenode
164         networks:
165             - lab5_ds_kafka_network
166         volumes:
167             - ./datanode2:/hadoop/dfs/data
168
169     datanode3:
170         image: bde2020/hadoop-datanode:latest
171         container_name: datanode3
172         environment:
```

```
173     - CORE_CONF_fs_defaultFS=hdfs://namenode:9000
174     depends_on:
175     - namenode
176     networks:
177     - lab5_ds_kafka_network
178     volumes:
179     - ./datanode3:/hadoop/dfs/data
180
181     networks:
182     lab5_ds_kafka_network:
183     external: true
184
185     volumes:
186     # Defines volumes for persistent storage of Namenode and Datanode data
187     namenode:
188     datanode:
189     datanode2:
190     datanode3:
```

After that, create the network **lab5_ds_kafka_network** by the cmd below in Docker:

```
1 docker network create lab5_ds_kafka_network
```

1.2 Previous LAB's Output

1 Check Output:

```
2 docker exec -it namenode bash -c 'hdfs dfs -ls /hdfs/environment_data'
```

```
3 docker exec -it namenode bash -c 'hdfs dfs -cat /hdfs/environment_data
```

```
/*.*.json'
```

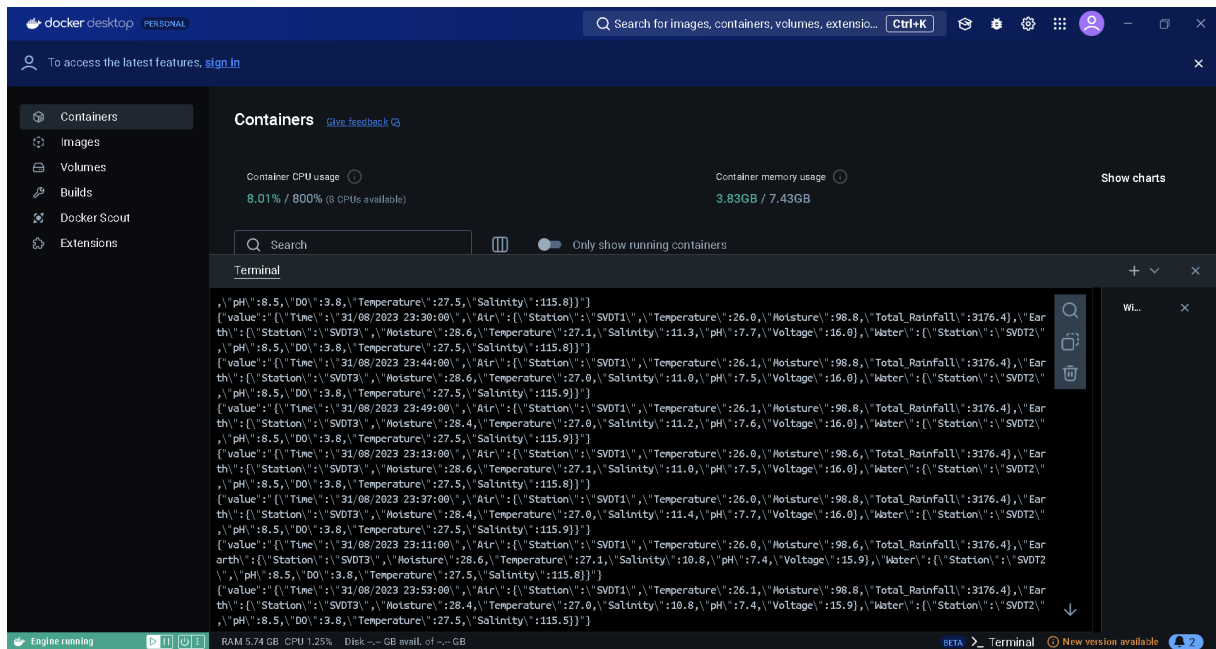


Figure 1: LAB4's Output

1.3 Loading Data

This section focuses on initializing the SparkSession and loading the environmental data stored in HDFS for distributed processing. The primary goal is to prepare the data for analysis of mango growth conditions.

- **SparkSession** serves as the entry point for PySpark applications, enabling distributed processing of large datasets. The session is named *"Distributed Computing for Mango Growth Evaluation"* for easy identification.
- The input data, stored as JSON files which contains information about air temperature, air moisture, and soil pH in HDFS at `hdfs://namenode:9000/hdfs/environment_data/*.json`.
- A schema is defined using **StructType** to parse the nested JSON structure accurately. It includes:
 - *Air*: Contains fields for temperature and moisture.
 - *Earth*: Contains the field for soil pH.
- The loaded data is parsed using `from_json`, and generating structured columns **Air_Temperature**, **Air_Moisture**, and **Earth_pH**, which are really vital for the subsequent analysis.

The parsed and structured dataset enables distributed computation while maintaining scalability and fault tolerance, leveraging the advantages of HDFS and Spark for efficient data handling.

```
1  from pyspark.sql import SparkSession
2  from pyspark.sql.functions import col, from_json, when, lit, avg, max, min
3  from pyspark.sql.types import StructType, StructField, DoubleType, StringType
4
5  # Initialize SparkSession
6  spark = SparkSession.builder \
7      .appName("Distributed Computing for Mango Growth Evaluation") \
8      .getOrCreate()
9
10 # HDFS Path
11 input_path = "hdfs://namenode:9000/hdfs/environment_data/*.json"
12 output_path = "hdfs://namenode:9000/hdfs/mango_data_distributed.txt"
13
14 # Define Schema for Nested JSON
15 schema = StructType([
16     StructField("Air", StructType([
17         StructField("Temperature", DoubleType()),
18         StructField("Moisture", DoubleType())
19     ])),
20     StructField("Earth", StructType([
21         StructField("pH", DoubleType())
22     ]))
23 ])
24
25 # Define Ideal Conditions
26 ideal_conditions = {
```

```
27     "Air_Temperature": (23.0, 28.0),
28     "Air_Moisture": (60.0, 80.0),
29     "Earth_pH": (5.5, 7.0)
30 }
31
32 def process_data():
33     # Load JSON Data and collect to driver
34     df = spark.read.json(input_path)
35
36     # Parse the "value" field
37     parsed_df = df.withColumn("parsed_value", from_json(col("value"), schema))
38
39     # Extract the parsed fields
40     data_df = parsed_df.select(
41         col("parsed_value.Air.Temperature").alias("Air_Temperature"),
42         col("parsed_value.Air.Moisture").alias("Air_Moisture"),
43         col("parsed_value.Earth.pH").alias("Earth_pH")
44     )
```

1.4 Bring data to compute (Centralized Computing)

This section explains the centralized computation approach, where the data is first collected from the distributed HDFS system and processed locally.

- **Data Collection:** All records from the distributed dataset are gathered to the driver node using the `collect()` method. This transforms the distributed data into a local collection, enabling computation to occur in a centralized environment.
- **Identifying Out-of-Ideal Conditions:** The script evaluates whether environmental parameters (`Air_Temperature`, `Air_Moisture`, `Earth_pH`) fall outside their ideal ranges:
 - `Air_Temperature`: Between 23°C and 28°C.
 - `Air_Moisture`: Between 60% and 80%.
 - `Earth_pH`: Between 5.5 and 7.0.

Each parameter is flagged with a 1 if it is out of range and 0 otherwise.

- **Real-Time Monitoring:** Using `show()` function outputs the `DataFrame` to monitor conditions in real-time, providing insights into the current state of the environment.
- **Time Analysis:** The total number of rows (representing minutes) where one or more parameters are out of range is counted. This is converted into hours and minutes for ease of understanding.
- **Statistical Analysis:** Using `agg()` function, the maximum and minimum values for each parameter are computed to assess the range and extremes of environmental conditions, offering valuable insights into the dataset's variability.

This method highlights the trade-off of centralized computing: while simple to implement, it may not scale efficiently for large datasets due to data transfer overheads.

```
1      # Collect all data to driver node
2      collected_data = data_df.collect()
3
4      # Convert collected data to a new local DataFrame
5      local_df = spark.createDataFrame(collected_data)
6
7      # Calculate Conditions Out of Ideal
8      out_of_ideal_df = local_df.withColumn(
9          "Air_Temperature_Out",
10         when((col("Air_Temperature") <
11             lit(ideal_conditions["Air_Temperature"][0])) |
12             (col("Air_Temperature") >
13                 lit(ideal_conditions["Air_Temperature"][1])), 1).otherwise(0)
14     ).withColumn(
15         "Air_Moisture_Out",
16         when((col("Air_Moisture") < lit(ideal_conditions["Air_Moisture"][0])) |
17             (col("Air_Moisture") > lit(ideal_conditions["Air_Moisture"][1])),
18             1).otherwise(0)
19     ).withColumn(
20         "Earth_pH_Out",
```

```
18     when((col("Earth_pH") < lit(ideal_conditions["Earth_pH"][0])) |
19           (col("Earth_pH") > lit(ideal_conditions["Earth_pH"][1])),
20           1).otherwise(0)
21 )
22
23 # Calculate time out of ideal conditions locally
24 out_of_ideal_count = out_of_ideal_df.filter(
25     (col("Air_Temperature_Out") + col("Air_Moisture_Out") +
26       col("Earth_pH_Out")) >= 1
27 ).count()
28
29 # Convert from minutes to hours, minutes
30 hours = out_of_ideal_count // 60
31 minutes = out_of_ideal_count % 60
32
33 # Calculate statistics locally
34 stats = local_df.agg(
35     max("Air_Temperature").alias("Max_Air_Temperature"),
36     min("Air_Temperature").alias("Min_Air_Temperature"),
37     max("Air_Moisture").alias("Max_Air_Moisture"),
38     min("Air_Moisture").alias("Min_Air_Moisture"),
39     max("Earth_pH").alias("Max_Earth_pH"),
40     min("Earth_pH").alias("Min_Earth_pH")
41 ).collect()[0]
```

1.5 Bring compute to data (Distributed Computing)

This section demonstrates the distributed computing approach, where the computation is brought to the data stored in HDFS, reducing data transfer overhead and leveraging distributed resources.

- **Identifying Out-of-Ideal Conditions:** Similar to the centralized approach, new columns are added to evaluate whether `Air_Temperature`, `Air_Moisture`, and `Earth_pH` fall outside their respective ideal ranges. This computation occurs in parallel across the distributed nodes, enhancing efficiency.
 - `Air_Temperature`: Ideal range is between 23°C and 28°C.
 - `Air_Moisture`: Ideal range is between 60% and 80%.
 - `Earth_pH`: Ideal range is between 5.5 and 7.0.
- **Real-Time Monitoring:** `show()` function is also utilized to display the `DataFrame`, allowing real-time observation of the conditions and offering insights into the current environmental state.
- **Time Analysis:** The script calculates the total time (in minutes) where one or more parameters are outside the ideal range. This is further converted to hours and minutes for better interpretability.
- **Statistical Summary:** The maximum and minimum values for each parameter are computed using Spark's `agg()` function. These calculations are distributed across the cluster, reducing processing time while handling large-scale data.

This approach effectively utilizes distributed computing capabilities to analyze environmental data at scale, minimizing the cost of moving large datasets.

```
1  # Calculate Conditions Out of Ideal
2  out_of_ideal_df = data_df.withColumn(
3      "Air_Temperature_Out",
4      when((col("Air_Temperature") < lit(ideal_conditions["Air_Temperature"][0])) |
5          (col("Air_Temperature") > lit(ideal_conditions["Air_Temperature"][1])),
6          1).otherwise(0)
7  ).withColumn(
8      "Air_Moisture_Out",
9      when((col("Air_Moisture") < lit(ideal_conditions["Air_Moisture"][0])) |
10         (col("Air_Moisture") > lit(ideal_conditions["Air_Moisture"][1])),
11         1).otherwise(0)
12 ).withColumn(
13     "Earth_pH_Out",
14     when((col("Earth_pH") < lit(ideal_conditions["Earth_pH"][0])) |
15         (col("Earth_pH") > lit(ideal_conditions["Earth_pH"][1])), 1).otherwise(0)
16 )
17
18 # Show the DataFrame
19 out_of_ideal_df.show()
20
21 # Calculate Average Time Out of Ideal Conditions
22 avg_out_of_ideal = out_of_ideal_df.filter((col("Air_Temperature_Out") +
23     col("Air_Moisture_Out") + col("Earth_pH_Out")) >= 1).count()
```

```
22  # Convert from minutes to hours, minutes, seconds
23  avg_out_of_ideal_hours = avg_out_of_ideal // 60
24  avg_out_of_ideal_minutes = avg_out_of_ideal % 60
25
26  # Calculate Max and Min for Each Field
27  stats_df = data_df.agg(
28      max("Air_Temperature").alias("Max_Air_Temperature"),
29      min("Air_Temperature").alias("Min_Air_Temperature"),
30      max("Air_Moisture").alias("Max_Air_Moisture"),
31      min("Air_Moisture").alias("Min_Air_Moisture"),
32      max("Earth_pH").alias("Max_Earth_pH"),
33      min("Earth_pH").alias("Min_Earth_pH")
34  ).collect()[0]
```

1.6 Store Output

The final step involves saving the processed results back to HDFS in a structured format for future use and further analysis.

- **Formatting Results:** The computed statistics and time metrics are formatted into a readable string that summarizes the analysis.
- **Storing Results:** The function `saveAsTextFile()` writes the formatted results as a single text file to the specified HDFS directory. This ensures the outputs are accessible and reusable for downstream processes.
- **Finalization:** `SparkSession` is terminated to free up resources, signifying the completion of the computation process.

This step ensures that all insights derived from the distributed computing process are preserved in a distributed storage system, enabling easy sharing and accessibility.

```
1      # Format results
2      results = (
3          f"Average Time Out of Ideal: {hours} hours {minutes} minutes\n"
4          f"Max Air Temperature: {stats['Max_Air_Temperature']}°C\n"
5          f"Min Air Temperature: {stats['Min_Air_Temperature']}°C\n"
6          f"Max Air Moisture: {stats['Max_Air_Moisture']}%\n"
7          f"Min Air Moisture: {stats['Min_Air_Moisture']}%\n"
8          f"Max Earth pH: {stats['Max_Earth_pH']}\n"
9          f"Min Earth pH: {stats['Min_Earth_pH']}\n"
10     )
11
12     # Write results to HDFS as a single text file
13     spark.sparkContext.parallelize([results]).saveAsTextFile(output_path)
14
15     # Process the data
16     process_data()
17
18     # Stop SparkSession
19     spark.stop()
```

1.7 Demonstration

1.7.1 Centralized Computing Method

```
1 Push code to Spark:
2   docker cp data_to_compute.py spark-master:/opt/spark/python/
   data_to_compute.py
3 Run code:
4   docker exec -it spark-master bash
5   spark-submit \
6     --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.3 \
7     --master spark://spark-master:7077 \
8     /opt/spark/python/data_to_compute.py
```

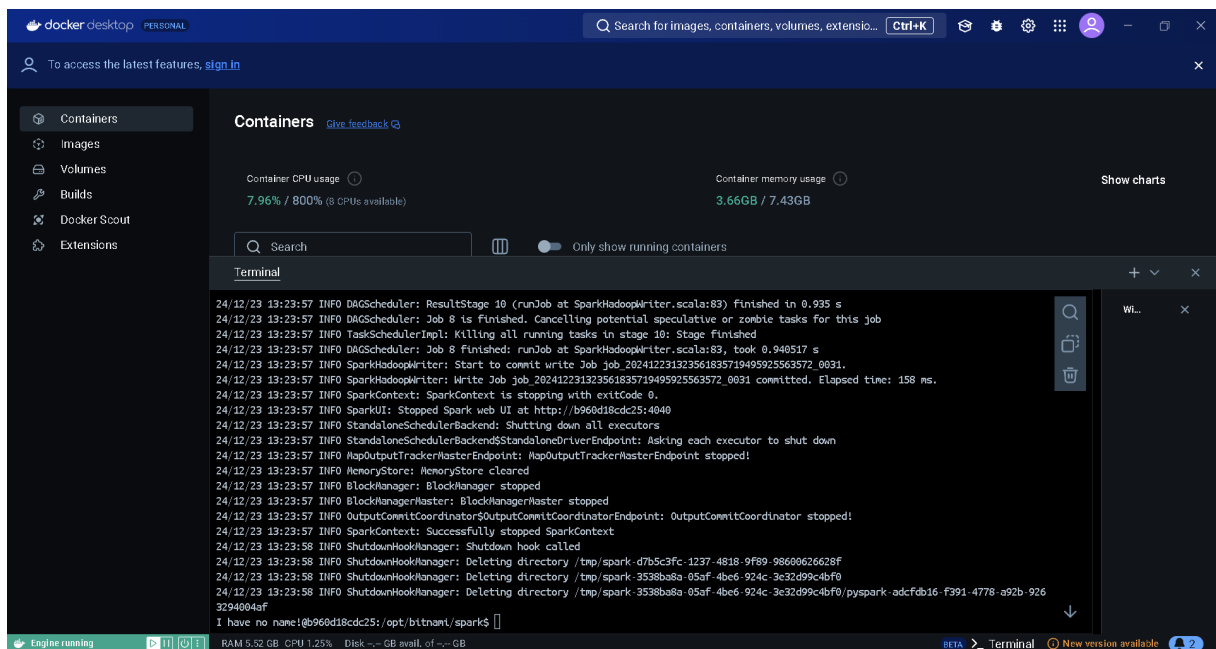


Figure 2: Finish Centralized Computing

```
1 Check Output:
2   docker exec -it namenode bash -c 'hdfs dfs -cat /hdfs/
   mango_data_centralized.txt/part-*
```

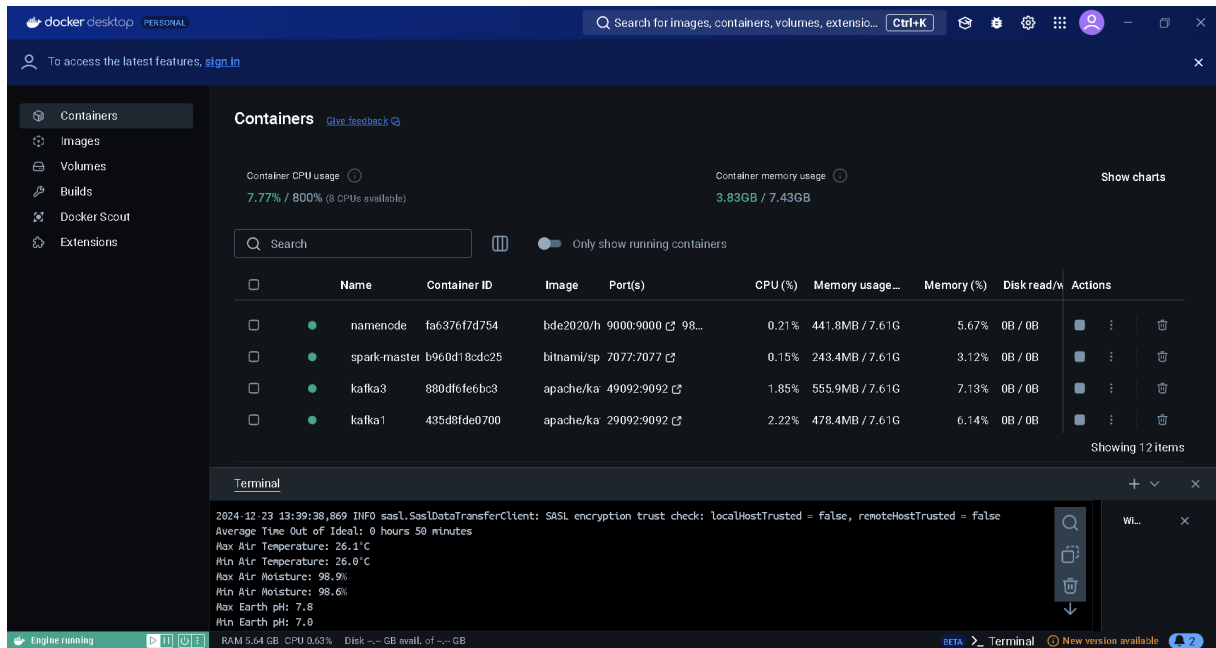



Figure 3: Output Centralized Computing

1.7.2 Distributed Computing Method

```
1 Push code to Spark:
2 docker cp compute_to_data.py spark-master:/opt/spark/python/
  compute_to_data.py
3 Run code:
4 docker exec -it spark-master bash
5 spark-submit \
6   --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.3 \
7   --master spark://spark-master:7077 \
8   /opt/spark/python/compute_to_data.py
```

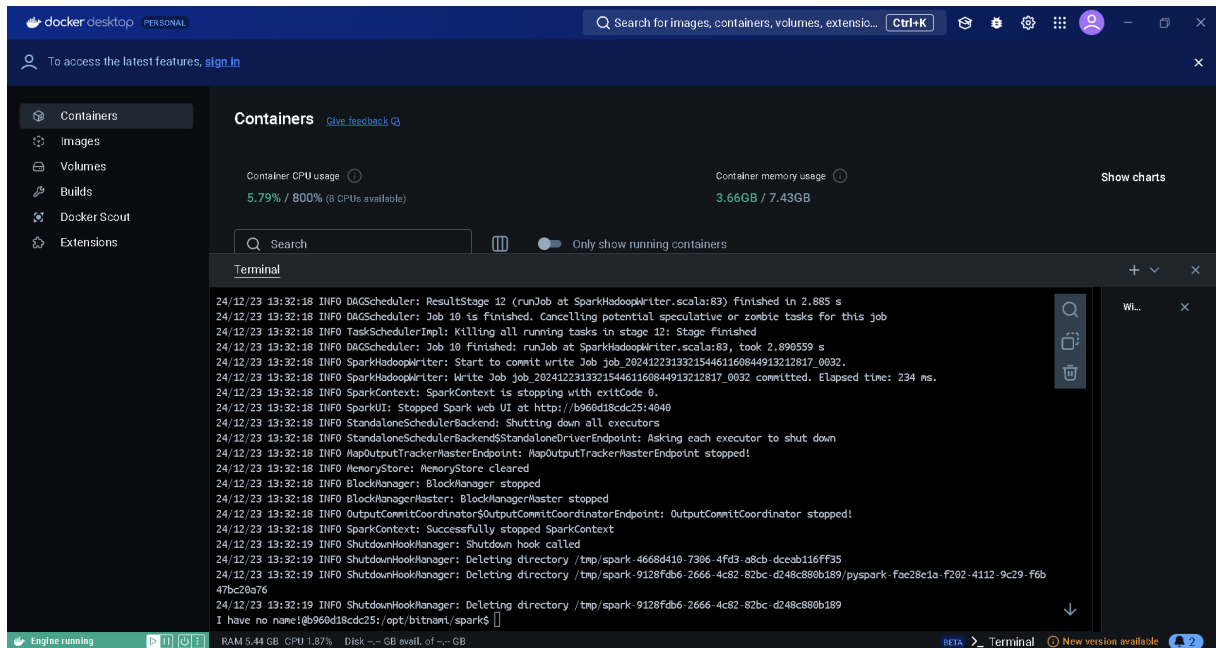


Figure 4: Finish Distributed Computing

1 Check Output:

```
2 docker exec -it namenode bash -c 'hdfs dfs -cat /hdfs/
  mango_data_distributed.txt/part-*
```

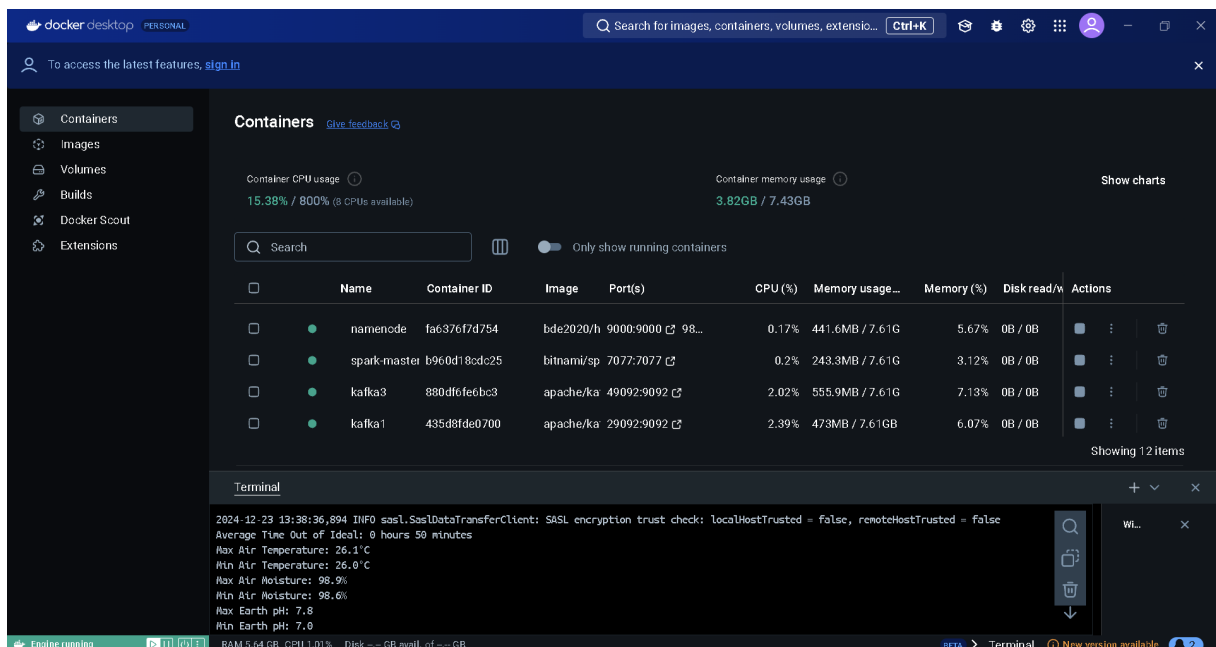


Figure 5: Output Distributed Computing

Observing the two output images generated by the two methods reveals that the results of both approaches are identical.